

cs-mail: Permissioned Communication One Relationship Request, One Conditional Charge

C-SQD

Draft 0.6 — September 2026

Abstract

cs-mail makes permission to communicate explicit. Accepted contacts write without protocol message charges. A sender without permission may submit one relationship request, accompanied by a communication and a conditional charge. Acceptance produces a full refund; expiry retains a processing amount; rejection also forfeits the collateral component. Forfeitures support a company-owned pool for later equal distributions among qualifying active members, after a disclosed corporate allocation. The recipient does not receive the money from the person they rejected. Native content remains end-to-end encrypted, and the system separates communication permission, financial obligations, and external payment execution.

1 Permission Before Repeated Interruption

An email address is easy to copy and difficult to protect. Ordinary mail lets someone attempt delivery repeatedly, leaving the recipient and their provider to sort wanted correspondence from unwanted contact. Filters can help, but they do not give the sender a clear account of whether an ongoing relationship exists.

cs-mail starts with that relationship. A recipient accepts a particular public identity, granting permission for future communication in that direction. Accepted correspondence carries no protocol message charge. The exceptional event is a request to establish permission where it is absent.

The sender makes one conditional payment for that request. The initial message gives the recipient context; permitted follow-ups remain part of the same request. The payment is not a purchase of attention, priority, or a reply. Its purpose is to make unwanted relationship requests costly while returning the charge when the recipient welcomes the relationship.

2 Alice Requests a Relationship with Bob

Bob has not accepted Alice. She sees a dollar-denominated quote with a processing component of 0.25 and collateral of 4.75, for a total charge of 5.00. She reviews the outcomes and decision window, authorizes the request, and supplies an encrypted initial message. The service admits it after confirming the required payment and durable delivery preparation. That admission starts one decision clock and creates one request for Bob to consider.

Outcome	The original charge	Future contact
Accept	All 5.00 is due back to Alice.	Alice gains standing permission to write to Bob.
Reject	0.25 is processing revenue; 4.75 becomes pending forfeiture.	Alice remains unaccepted and must wait until another request is eligible.
Expire	0.25 is processing revenue; 4.75 is due back to Alice.	No standing permission is granted.

Bob may allow follow-up messages while this request is open. If he does, Alice can add context without another charge, another request prompt, or a later deadline. If he does not, the service refuses the follow-up. A sender cannot evade that choice by paying again against the same open request.

When Bob accepts, the service records permission and the refund obligation together. The payment provider may confirm the refund afterward, so Alice’s statement distinguishes a refund due from a refund completed. Bob’s acceptance is directed: it does not automatically accept Bob as a sender to Alice or grant permission to every other identity Alice controls.

3 The Request Has a Beginning and an End

Preparing a request is different from admitting it. An abandoned or cancelled request that has not been admitted receives a full void or refund and produces no recipient solicitation. If payment succeeds after cancellation, the service still refunds it rather than reviving the request. Closing an application alone is not proof that a payment was cancelled.

After admission, ordinary delivery retries do not reopen the sender’s cancellation option. The request remains subject to the recipient’s decision or its fixed expiry. This prevents cancellation from becoming a way to escape rejection. A message’s own delivery-validity limit remains separate from that decision clock.

Rejection and blocking express different intentions. Rejection declines this request but may permit a later eligible one. Blocking prohibits contact until an authenticated unblock, including communication through a previously issued lane. Blocking an open request uses its rejection allocation; blocking before admission returns the charge in full. Blocking after finalization does not reopen money.

The service records when an authenticated decision arrives. A timely acceptance must not lose to an expiry worker merely because processing was delayed. A late acceptance can grant future permission, but it cannot change an already closed request’s financial outcome. These rules make deadline behavior explainable to both parties.

4 Where Forfeited Collateral Goes

Let C be the processing component and S the refundable collateral component. The request charge is $Q = C + S$. Acceptance returns both; expiry returns S ; rejection makes S a pending forfeiture. No part of this rule gives the rejecting recipient a claim on that sender’s collateral.

Instead, eligible forfeitures support a C-SQD-owned pool. “Eligible” matters: money that may still be needed for refunds, disputes, or required reserves cannot fund distributions. Separate accounting protects these obligations even when the underlying cash is held through the same financial provider.

The proposed corporate allocation is 3% of newly eligible forfeitures, assessed once under the rate disclosed for each request. The remaining 97% enters the member-restricted portion. For

example, 10,000 of newly eligible forfeitures would yield 300 of corporate revenue and 9,700 for members. This is not an extra charge to the sender. Previously assessed member funds are not charged another 3% each quarter.

C-SQD cannot spend the member-restricted portion on operations or payout fees. Those costs come from its processing revenue and corporate allocation. The exact launch rate remains a policy choice; the separation between corporate revenue, restricted funds, and outstanding obligations is part of the model.

5 A Quarterly Benefit for Active Members

The pool benefits qualifying active members generally, rather than only those who receive unsolicited requests. Each qualifying member receives an equal share. Sending more requests, rejecting more people, or choosing more expensive collateral does not increase that share. Multiple mailboxes or devices do not create multiple memberships.

At a published quarterly cutoff, C-SQD fixes the eligible membership and funding records. It then finalizes equal allocations and records the amount owed to each member. A failed payment or an amount below a disclosed payment threshold remains that member's payable; it does not return to the next quarter's equal split. Unallocated fractional-unit remainder stays in the restricted pool. If no members qualify, unallocated funds carry forward.

The proposed activity rule is modest: membership for at least 30 days and intentional authenticated use on two distinct days during the quarter. Reading can count; background polling alone does not. These are candidate defaults, along with verification and opt-in requirements. Organization eligibility and the remaining launch parameters are identified in the deployment profile rather than treated as settled facts here.

This program is distinct from a wallet. Members cannot spend a pending allocation on another request, transfer it to someone else, or deposit generally spendable funds. External providers execute charges, refunds, and payments of finalized obligations. The ledger records what the service owes and what has been confirmed.

6 Persistence Without Repeated Charges Inside a Request

One request per directed identity pair does not by itself stop someone from returning after every expiry or rotating addresses. New requests therefore remain subject to recipient policy and principal-scoped history. Initial admission advances that history; an allowed follow-up does not. Acceptance of a public identity bypasses the request gate for that identity without erasing the history that applies to its aliases.

The deployment proposal uses a 30-day decision window and cooldowns after expiry or rejection, with recipient controls and explicit invitations. The final durations belong in published, versioned policy. Blocking remains effective regardless of willingness to pay.

Earlier drafts also proposed a refundable persistence reserve L_k . Its future use remains unresolved. The revised base model does not enable a nonzero reserve; retaining one would require explicit charge and release rules. It could attach only to a new request and would never fund the pool or corporate revenue.

7 Permission for Organizations: Express Lanes

Some wanted correspondence is narrower than standing personal acceptance. A member may want a bank's fraud alerts or an airline's itinerary notices without inviting unlimited communication from that organization. An express lane is a recipient-signed, revocable grant bounded by sender, declared purpose, count or rate, and time.

A valid lane authorizes messages without a relationship-request charge. Its messages do not contribute to forfeitures or create a request solicitation. Granting, renewing, or revoking a lane does not settle a separate pending request. A block overrides the lane. If lane authority ends, the sender must satisfy the remaining permission rules, including any existing request's follow-up policy; another charge is never automatic.

Native lanes use cryptographic identity and scoped references. Legacy lanes use domain-authentication evidence at a gateway and are visibly subject to that route's privacy limits. A signed purpose declaration says who made the claim; neither a signature nor domain authentication proves the meaning of encrypted content. The companion express-lane documents explain issuance and lifetime.

8 Privacy Covers Relationships as Well as Content

The words in a message are only part of its sensitivity. Who tried to contact whom, who accepted or rejected, and how financial outcomes followed can reveal personal and professional relationships. Native content is therefore encrypted at the endpoints, while the surrounding services receive only the references and facts needed for their roles.

Public protocol identities are the personas people address and accept. A private principal connects control, recovery, and anti-abuse continuity inside its identity authority. A distribution-program member is a third concept: a verified subject eligible for one share. These identities should not collapse into one global identifier passed through every service.

Relationship and financial systems use separately scoped opaque handles. A publicly testable hash of two addresses would not provide that protection because an observer could guess the addresses. Access to ciphertext storage should not grant access to identity evidence; membership verification should not expose a contact graph. Separate keys, credentials, and audited translation boundaries make those distinctions operational.

The envelope may expose a minimal authenticated purpose, origin, validity time, and permission or request reference needed for delivery. It does not expose subjects, bodies, attachments, legal identity, or unrelated financial activity. The external payment provider receives required request/payment references and amounts, not the message history within the request. Recipient follow-up rules are admission policy; ranking and notification preferences are presentation policy. Neither silently rewrites financial terms.

Privacy also depends on time. A pending request needs enough linkage for delivery, decisions, and refunds. After its operational and required retention periods, unnecessary linkage can be deleted while minimal financial obligations remain auditable. An accepted relationship may persist without preserving every message that formed it. Unpaid member allocations retain the evidence needed for payment, not a permanent archive of correspondence.

Encryption is not erasure if the operator retains usable keys and identity mappings. Deletion must cover replicas, backups, logs, and derived indexes, with explicit treatment of required records. A consent-locked historical archive may use independent participant key shares, but live permission and payment must not depend on both participants keeping those shares. No system can erase

copies independently retained by a sender or recipient.

9 The Ledger Makes the Promise Testable

The canonical ledger uses append-only double-entry accounting. It distinguishes request amounts, refunds due, pending forfeitures, corporate allocations, restricted funds, and member payables. A recipient decision commits its financial obligations and retryable external work together. Payment confirmation arrives as a separate fact. Duplicate commands and provider notifications must not repeat the economic effect.

The resulting audit can answer concrete questions: was this request charged once; did acceptance produce the promised refund; was a forfeiture eligible; was its corporate share assessed once; did every qualifying member receive an equal allocation; and does every unpaid allocation remain recorded? These checks do not require publishing the underlying correspondence graph.

The incentives still deserve scrutiny. The operator earns processing revenue on unsuccessful requests and a share of eligible forfeitures. Equal distributions remove the direct rejection bounty but do not eliminate duplicate-member farming, fabricated activity, collusion, or manipulation of deadlines. Visible acceptance, signed receipts, pricing controls, scoped review, and aggregate reconciliation remain necessary parts of the design.

10 Building and Operating the First Deployment

C-SQD can begin as one operator with a deterministic request kernel, a durable ledger, encrypted delivery, and external payment adapters. The desktop pilot can simulate charges and quarterly distributions without real redeemable value. That allows the request and failure stories to be tested before members depend on live payments.

Real charges require completed payment arrangements, maturity and loss policies, member eligibility rules, reconciliation, and orderly closure provisions. Company ownership, pooling, and an external provider are architectural choices to evaluate, not a conclusion that the service is exempt from financial regulation. Future federation also needs explicit allocation of pool governance, membership uniqueness, refund obligations, clearing, and losses.

The *Protocol Specification* defines the common request rules. The *Deployment and Migration Profile* defines the C-SQD financial program and identifies open launch parameters. The architecture and desktop plan translate that model into staged implementation. Together they preserve the central experience: permission is explicit, an open request carries one charge, and ordinary communication within an accepted relationship carries none.